


14013204-3 - PARALLEL COMPUTING



Chapter 2

Parallel Software And Performance

Roadmap

- Parallel software
- Measuring Performance
- Sources of performance loss
- Performance trade-offs



PARALLEL SOFTWARE

The burden is on software

- Parallel hardware has arrived.
 - Virtually all desktop and server systems use multicore processors. Nowadays even mobile phones and tablets make use of multicore processors.
- The situation for parallel software is in flux.
 - Most system software makes some use of parallelism, and many widely used application programs (e.g., Excel, Photoshop, Chrome) can also use multiple cores.
 - However, there are still many programs that can only make use of a single core, and there are many programmers with no experience writing parallel programs.

The burden is on software

- Can no longer rely on Hardware and compilers to provide a steady increase in application performance.
- Software developers must learn to write applications that exploit shared- and distributed-memory architectures and MIMD and SIMD systems.
- we'll take a look at some of the issues involved in writing software for parallel systems.
- From now on...
 - In shared memory programs:
 - Start a single process and fork threads.
 - Threads carry out tasks.
 - In distributed memory programs:
 - Start multiple processes.
 - Processes carry out tasks.

SPMD – single program multiple data

- We will mainly focus on **SPMD** programs.
- A SPMD program consists of a single executable that can behave as if it were multiple different programs through the use of conditional branches.
- Implement both **Task parallel** and **Data parallel** programs.

```
if (I'm thread process i)
    do this;
else
    do that;
```



Coordinating the processes/threads

1. Divide the work among the processes/threads

- (a) so each process/thread gets roughly the same amount of work.
- (b) and communication is minimized.

2. Assign (allocate) the work to processes/threads.

3. Arrange for the processes/threads to synchronize.

4. Arrange for communication among processes/threads.

- These last two problems are often **interrelated**. For example, in distributed-memory programs, we often implicitly synchronize the processes by communicating among them, and in shared-memory programs, we often communicate among the threads by synchronizing them.
- **Parallelization:** The process of converting a serial program or algorithm into a parallel program.
- **Load Balancing:** The process of dividing the work among the processes/threads so that each process/thread gets roughly the same amount of work.

```
double x[n], y[n];  
...  
for (i = 0; i < n; i++)  
    x[i] += y[i];
```

Shared Memory

- In shared-memory programs, variables can be **shared** or **private**.
 - Shared variables can be read or written by any thread.
 - Private variables can ordinarily only be accessed by one thread.
- Communication among the threads is usually done through shared variables, so communication is **implicit** rather than **explicit**.

Shared Memory

- **Dynamic threads**

- Master thread waits for work, forks new threads, and when threads are done, they terminate.
- Efficient use of resources, but thread creation and termination is time consuming.

- **Static threads**

- Pool of threads created and are allocated work, but do not terminate until cleanup.
- Better performance, but potential waste of system resources.

Nondeterminism

- In any MIMD system in which the processors execute asynchronously it is likely that there will be **nondeterminism**.
- A computation is nondeterministic if a given input can result in different outputs.
- If multiple threads are executing independently, the relative rate at which they'll complete statements **varies from run to run**, and hence the **results of the program may be different from run to run**.

```
...
printf ( "Thread %d > my_val = %d\n" ,
        my_rank , my_x );
...
```

Thread 1 > my_val = 19

Thread 0 > my_val = 7

Thread 0 > my_val = 7

Thread 1 > my_val = 19

Nondeterminism

```
my_val = Compute_val ( my_rank ) ;  
x += my_val ;
```

Time	Core 0	Core 1
0	Finish assignment to my_val	In call to Compute_val
1	Load x = 0 into register	Finish assignment to my_val
2	Load my_val = 7 into register	Load x = 0 into register
3	Add my_val = 7 to x	Load my_val = 19 into register
4	Store x = 7	Add my_val to x
5	Start other work	Store x = 19

- The nondeterminism here is a result of the fact that two threads are attempting to more or less simultaneously update the memory location **x**.
- When threads or processes attempt to simultaneously access a shared resource, and the accesses can result in an error, we often say the program has a **race condition**, because the threads or processes are in a “**race**” to carry out an operation.
- That is, **the outcome of the computation depends on which thread wins the race**.

Nondeterminism

- **Critical section:** is a block of code that can only be executed by one thread at a time.
- It's usually our job as programmers to ensure **Mutually Exclusive** access to a critical section.
 - In other words, we need to ensure that if one thread is executing the code in the critical section, then the other threads are excluded.
- **Mutual exclusion lock (mutex, or simply lock):** The most commonly used mechanism for ensuring mutual exclusion. A mutex is a special type of object that has support in the underlying hardware.
 - No pre-determined order on the threads.
 - The code in the critical section is sequential.

```
my_val = Compute_val ( my_rank ) ;  
Lock(&add_my_val_lock ) ;  
x += my_val ;  
Unlock(&add_my_val_lock ) ;
```

busy-waiting

- There are alternatives to mutexes. In busy-waiting, a thread **enters a loop, whose sole purpose is to test a condition.**
- In our example, suppose there is a shared variable `ok_for_1` that has been initialized to false. Then something like the following code can ensure that thread 1 won't update `x` until after thread 0 has updated it:

```
my_val = Compute_val ( my_rank ) ;  
if ( my_rank == 1 )  
    while ( ! ok_for_1 ) ; /* Busy-wait loop */  
    x += my_val ; /* Critical section */  
if ( my_rank == 0 )  
    ok_for_1 = true ; /* Let thread 1 update x */
```


Distributed Memory

- In distributed-memory programs, the cores can directly access only their own, private memories.
- There are several APIs that are used. The most widely used is **message-passing**.
- distributed-memory programs are usually executed by starting multiple processes rather than multiple threads.
- A message-passing API provides (at a minimum) a **send** and a **receive** function. Processes typically identify each other by ranks in the range $0, 1, \dots, p - 1$, where p is the number of processes.

Message-Passing

```
char message [ 1 0 0 ] ;  
...  
my_rank = Get_rank ( ) ;  
if ( my_rank == 1) {  
    sprintf ( message , "Greetings from process 1" ) ;  
    Send ( message , MSG_CHAR , 100 , 0 ) ;  
} else if ( my_rank == 0) {  
    Receive ( message , MSG_CHAR , 100 , 1 ) ;  
    printf ( "Process 0 > Received: %s\n" , message ) ;  
}
```

sprintf in C is a library function used for formatted output to a string. It works similarly to printf, which prints to the console, but sprintf stores the output in a character buffer (string) specified as its first argument.



Input and Output

- We'll be making these assumptions and following these rules when our parallel programs need to do I/O:
 - In distributed memory programs, only process 0 will access *stdin*. In shared memory programs, only the master thread or thread 0 will access *stdin*.
 - In both distributed memory and shared memory programs all the processes/threads can access *stdout* and *stderr*.
 - However, because of the nondeterministic order of output to *stdout*, in most cases only a single process/thread will be used for all output to *stdout*.
 - The principal exception will be output for debugging a program. In this situation, we'll often have multiple processes/threads writing to *stdout*.
 - Debug output should always include the rank or ID of the process/thread that's generating the output.

Input and Output


- Only a single process/thread will attempt to access any single file other than *stdin*, *stdout*, or *stderr*.
- So, for example, each process/thread can open its own, private file for reading or writing, but no two processes/threads will open the same file.



PERFORMANCE

Measuring Performance

- Basic measure is *execution time*



Measuring Performance



- Performance is usually **main aim** of parallel computing
 - Some problems take too long on single processor.
- Usually, the best we can hope to do is to equally divide the work among the cores, with no additional work for the cores.
- If we succeed in doing this, and we run our parallel program with **p cores**, one thread or process on each core, then our **parallel program** will run **p times faster** than the **serial program**.
- This is often **difficult to achieve**, e.g.
 - Splitting work, Data dependencies, Load balancing, False sharing, etc.
- Let's define performance formally.

Measuring Performance

- If we call the serial **run-time** T_{serial} and our parallel run-time T_{parallel} , then the best we can hope for is

$$T_{\text{parallel}} = \frac{T_{\text{serial}}}{p}$$

- When this happens, we say that our parallel program has **linear speedup**.
- In practice, we usually **don't get perfect linear speedup**, because the use of multiple processes/threads almost invariably introduces some **overhead** (*(The amount of time required to coordinate parallel tasks, as opposed to do useful work)*).
 - It's likely that the overheads will increase as we increase the number of processes or threads.

Sequential Code / Parallel Code

```
int sum = 0;
for(int i = 0; i < size; i++) {
    sum += array[i];
    //sum=sum+array[i];
}
```

```
int numThreads = 2; //Assume one thread per core, and 2
cores
int sum = 0;
int i = 0;
int middleSum[numThreads];
int threadSetSize = size/numThreads
//Each thread will execute this code with a different
threadID
for( i = threadID*threadSetSize; i <
(threadID+1)*threadSetSize; i++)
{
    middleSum[threadID] += array[i];
}
waitForAllThreads(); //Wait for all threads
//Only thread 0 will execute this code
if (threadID==0) {
    for(i = 0; i < numThreads; i++) {
        sum += middleSum[i];}}
```

Measuring Performance


- Number of cores = p
- Serial run-time = T_{serial}
- Parallel run-time = T_{parallel}

linear speedup

$$T_{\text{parallel}} = T_{\text{serial}} / p$$

Measuring Performance

- Execution time is not enough
- How do we relate different times?
 - Maybe we know one time is better than another
 - But how much better? What is the relationship?



Speedup of a parallel program

- **Speedup** measures increase in running time due to parallelism.
- Execution time of sequential program divided by execution time of parallel program (for the same problem).

$$S = \frac{T_{\text{serial}}}{T_{\text{parallel}}}$$


- $T_{\text{serial}}(s)$ = Sequential time
- $T_{\text{parallel}}(p)$ = Parallel time on P processors
- E.g., T_2 is parallel time on 2 processors
- Linear speedup has $S = p$. Ideally, **speedup** is linear (perfect). E.g.

- $\frac{T_s}{T_2} = 2$
- $\frac{T_s}{T_4} = 4$


- But **this rarely happens**. There are several sources of performance loss. Will be discussed later.

Example Speedup Graph

- **super-linear speedup: $S > p$**
- **linear speedup: $S = p$**
- **sub-linear speedup: $S < p$**



Example Speedup Graph



Speedup

- T_s is time of sequential program
- NOT parallel program with 1 process/thread.
 - Parallel program has extra **overheads**
 - Example: array sum results present only parallel program with different number of threads.
 - **Sequential program:** 0.52 seconds (T_s)
 - **Parallel program, 1 thread:** 0.86 seconds (T_1)

Speedup

- Parallel program, 1 thread is $T_1 \neq T_s$
 - **Absolute speedup** = $T_s / T_2 = 0.52 / 0.45 = 1.16$
 - **Relative speedup** = $T_1 / T_2 = 0.86 / 0.45 = 1.91$

Efficiency of a parallel program

linear speedup has $S = p$, which is unusual.

- As p increases, we expect S to become a smaller and smaller fraction of the ideal, linear speedup p .
- Another way of saying this is that S/p will probably get smaller and smaller as p increases.
- This value, S/p , is sometimes called the **efficiency** of the parallel program. If we substitute the formula for S , we see that the efficiency is

$$E = \frac{S}{p} = \frac{\frac{T_{\text{serial}}}{T_{\text{parallel}}}}{p} = \frac{T_{\text{serial}}}{p \cdot T_{\text{parallel}}}$$

Efficiency of a parallel program

- If the serial run-time has been taken on the same type of core that the parallel system is using, we can think of **efficiency** as **the average utilization of the parallel cores on solving the problem**.
- That is, the efficiency can be thought of as **the fraction of the parallel run-time that's spent, on average, by each core working on solving the original problem**. ((To measure how effectively each processor used)).
- The remainder of the parallel run-time is the **parallel overhead**.
- If **efficiency = 1**, then **linear speedup**
- If **efficiency < 1**
 - Processors **not fully used** (to solve problem).
 - This is the normal case.
 - Many reasons for performance loss.
 - Also, efficiency falls as processors increased.

Effect of overhead

- For example, suppose we have $T_{\text{serial}} = 24 \text{ ms}$, $p = 8$, and $T_{\text{parallel}} = 4 \text{ ms}$. Then,

$$E = \frac{24}{8 \cdot 4} = \frac{3}{4},$$

- On average, each process/thread spends $3/4 \cdot 4 = 3 \text{ ms}$ on solving the original problem, and $4 - 3 = 1 \text{ ms}$ in **parallel overhead**.
- Many parallel programs are developed by explicitly dividing the work of the serial program among the processes/threads and adding in the necessary “parallel overhead,” such as mutual exclusion or communication.
- Therefore if T_{overhead} denotes this parallel overhead, it’s often the case that:

$$T_{\text{parallel}} = T_{\text{serial}} / p + T_{\text{overhead}}$$

Speedups and efficiencies of a parallel program

- We've already seen that T_{parallel} , S , and E , depend on p , the number of processes or threads.



p	1	2	4	8	16
S	1.0	1.9	3.6	6.5	10.8
$E = S/p$	1.0	0.95	0.90	0.81	0.68

Speedups and efficiencies of parallel program on different problem sizes

- We also need to keep in mind that T_{parallel} , S , E , and T_{serial} all depend on **the problem size**.
- For example, if we **halve and double** the problem size of the program, whose speedups are in the previous slide, we get the speedups and efficiencies shown:
 - when we increase the problem size, the speedups and the efficiencies increase, while they decrease when we decrease the problem size.
 - This behavior is quite common, because in many parallel programs, as the problem size is increased but the number of processes/threads is fixed, the parallel overhead T_{overhead} grows much more slowly than the time spent in solving the original problem.
 - There's more **computation** work for the processes/threads to do, so the relative amount of **coordination** time should be less.

	p	1	2	4	8	16
Half	S	1.0	1.9	3.1	4.8	6.2
	E	1.0	0.95	0.78	0.60	0.39
Original	S	1.0	1.9	3.6	6.5	10.8
	E	1.0	0.95	0.90	0.81	0.68
Double	S	1.0	1.9	3.9	7.5	14.2
	E	1.0	0.95	0.98	0.94	0.89

Speedup and Efficiency



Problem Size

- **Fixed-size speedup**

- Use same problem size for all processor counts.
- Problem: If large range of processors, then small problem size that fits in memory of one processor may be too small for, e.g., 100000 processors = high overhead = reduced speedup
- Solution: scaled speedup

- **Scaled speedup**

- Increase problem size with number of processors.
- Not straightforward. E.g., how is larger problem size affected by memory and interconnect architecture?

Taking Timings


- What is time?
- Start to finish?
- A program segment of interest?
- CPU time?
- Wall clock time?



Taking Timings

- We are interested in the time spent doing specific computation, not the entire execution of the parallel program.
- To measure the **wall clock time** for specific computation:

```
double start, finish;  
. . .  
start = Get_current_time();  
/* Code that we want to time */  
. . .  
finish = Get_current_time();  
printf("The elapsed time = %e seconds\n", finish-start);
```



- **Get current time()** is a function that's supposed to return the number of seconds that have elapsed since some fixed time in the past.

Taking Timings

- we're usually interested in is a single time: **the time that has elapsed from when the first process/thread began execution of the code to the time the last process/thread finished execution of the code.**

```
shared double global_elapsed;
private double my_start, my_finish, my_elapsed;
. .
/* Synchronize all processes/threads */
Barrier();
my_start = Get_current_time();

/* Code that we want to time */
. .

my_finish = Get_current_time();
my_elapsed = my_finish - my_start;

/* Find the max across all processes/threads */
global_elapsed = Global_max(my_elapsed);
if (my_rank == 0)
    printf("The elapsed time = %e seconds\n", global_elapsed);
```

Pitfalls (beware)

- **Do not compare speedup on different processors.**
- E.g.
 - Sequential on Pentium 4, $T_s = 1$ second
 - Parallel on 8x Pentium 1, $T_8 = 10$ seconds
 - Speedup = $T_s/T_8 = 0.1$
 - Efficiency = Speedup/P = $0.1/8 = 0.0125$
- **These results are meaningless**
 - Must use same hardware

Pitfalls (beware)

■ Cold starts

- First run of a program usually slow
- Page table misses (virtual memory)
- Cache misses
- Second, third, ++ runs are faster
- This means, **warm up memory/cache before you start measuring performance.**

Pitfalls (beware)

- Related issue: single-run performance
 - Don't calculate speedup based on one run, as runtimes vary
 - Need to **take average of several runs**, e.g. report either a **mean** or a **median** run-time:
 - Mean: add up all the data, and then divide this total by the number of values in the data.
 - Median: put the values in order, then find the middle value.
 - Example: Find the mean, and median for the following list of values:
13, 18, 13, 14, 13, 16, 14, 21, 13
 1. **Mean**: $(13 + 18 + 13 + 14 + 13 + 16 + 14 + 21 + 13) \div 9 = 15$
 2. **Median**: Rewrite the list in order: 13, 13, 13, 13, **14**, 14, 16, 18, 21

Pitfalls (beware)

■ I/O Accesses


- **Network, hard disk accesses take a lot of time**
- If program is regularly making such accesses, **performance is difficult to measure**, because **execution times vary** a lot.
- E.g.
 - Sequential: 6s
 - Parallel 2 threads:
 - 1st run: 5s (**cold start**)
 - 2nd run: 2s (**warm up**)
 - 3rd run: 4s (**another program accessing disk**)
- As a practical matter, since our programs won't be designed for high-performance I/O, we'll usually not include I/O in our reported run-times.

Sources of Performance Loss

- Overhead
- Contention for resources
- Idle time
- Non-parallelizable computation

Overhead

- Any performance cost of parallelizing a sequential program
- Any cost that is incurred in the parallel solution but not in the sequential one.
- **Process + thread creation/destruction**
 - Sequential program only creates one process.
 - Parallel program creates **at least one process (maybe more)**, and **at least one thread (usually more)**. Also have to destroy them at the end.
 - **This results in reduced performance.**
 - with array sum: $T_s = 0.52$, $T_1 = 0.86$



Overhead

- Four main sources

1. **Communication**
2. **Synchronization**
3. **Computation**
4. **Memory**

Overhead

■ **Communication**


- **Threads and processes** communicate.
- Distributed-memory programs will almost always need to transmit data across the network, which is usually much slower than local memory access.
- False sharing in shared memory systems is also communication. (**unwanted**)
- Sequential program doesn't use communication (1 process).

■ **Synchronization**

- Usually **causes one thread to wait for another**.
- Shared-memory programs will almost always have **critical sections**.
 - Which will require that we use some **mutual exclusion** mechanism, such as a **mutex**. The calls to the mutex functions are the overhead.
 - The use of the such synchronization mechanisms forces the parallel program to serialize execution of the critical section.
 - Sequential program doesn't use synchronization (1 process)
- The overheads **will increase as we increase the number of processes or threads**. For example, more threads will probably mean more threads need to access a critical section, and more processes will probably mean more data needs to be transmitted across the network.

Synchronization Cost Example

- **Barrier overhead**
 - **Gray: threads working**
 - **Red: threads waiting**
 - Load balance
 - **Black: communication:**
 - message all threads to resume



Overhead

- **Computation**

- Parallel solution usually **does extra computation**.
- Example from array sum: *middle sums*.

- **Memory**

- **Parallel solution** may use **more memory**.
- Array sum example: **padding memory** to remove false sharing.
- This actually improves performance but increases **cost of memory**. May not always be possible if memory limit reached by program.

Contention

- **Competition for a shared resource**
 - E.g., bus in the computer's architecture.
- **Parallel program increases contention**
 - More processors/threads accessing shared resources.
 - Can affect processors that are working on different areas of the program.
 - E.g., imagine **two processors causing contention on a bus due to false sharing**.
 - A third processor's **memory request latency is increased**, even if it is not involved in the false sharing.

Idle Time

- Ideally, processors working all the time.
- Not usually the case
 - Waiting for data from memory
 - Load imbalance (not enough work)

Non-Parallelizable Code

- **Some code cannot be parallelized**
 - Code with **dependencies** is one example
 - This code **has to be executed sequentially**
 - This code **limits the benefit of parallelization**

Amdahl's Law



- Unless virtually all of a serial program is parallelized, the possible speedup is going to be very limited — regardless of the number of cores available.
- Example:** Suppose that we can parallelize 90% of a serial program.
- Suppose that the speedup of the parallelized part is perfect, which means that the speedup is always p regardless of the number of p .
 - If the sequential run-time is $T_{\text{serial}} = 20$ seconds
 - Runtime of parallelizable part is $0.9 \times T_{\text{serial}} / p = 18 / p$
 - Runtime of “unparallelizable” part is $0.1 \times T_{\text{serial}} = 2$
 - Overall parallel run-time is:** $T_{\text{parallel}} = 0.9 \times T_{\text{serial}} / p + 0.1 \times T_{\text{serial}} = 18 / p + 2$
 - Speed up:**
$$S = \frac{T_{\text{serial}}}{0.9 \times (T_{\text{serial}} / p) + 0.1 \times T_{\text{serial}}} = \frac{20}{(18 / p) + 2}$$

Amdahl's Law




- Now as p gets larger and larger $0.9 \times \frac{T_{sequential}}{p} = \frac{18}{p}$ gets closer and closer to 0.
- So the total parallel runtime can't be smaller than $0.1 \times T_{sequential} = 2$
- That is, the denominator in S can't be smaller than $0.1 \times T_{sequential} = 2$
The fraction S must therefore be smaller than: $S \leq \frac{T_{sequential}}{(0.1 \times T_{sequential})} = \frac{20}{2} = 10$
that is $S \leq 10$
- This is saying that even though we've done a perfect job in parallelizing 90% of the program, and even if we have, say, 1000 cores, we'll never get a speedup better than 10.

Amdahl's Law



- More generally, if a **fraction r** of our **sequential program remains unparallelized**, then Amdahl's law says **we can't get a speedup better than $\frac{1}{r}$** .
In our example, $r = 1 - 0.9 = \frac{1}{10}$ so we couldn't get a **speedup better than 10**.
- Thus, even if r is quite small, e.g. say $\frac{1}{100}$, and we have a system with **thousands** of cores, we can't possibly get a speedup better than 100.

Amdahl's Law



Amdahl's Law

- Shall we give up? **No**
- There are **several reasons** not to be too worried by Amdahl's law:
 1. **it doesn't take into consideration the problem size.**
 - Often when we have more processors, we also increase the problem size.
 - E.g. **increase amount of data to be computed.**
 - This reduces sequential part of the total problem, so we can get more speedup.
 2. There are **thousands of programs** used by scientists and engineers that routinely **obtain huge speedups on large distributed-memory systems.**
 3. In many cases, **obtaining a speedup of 5 or 10 is more than adequate**, especially if the effort involved in developing the parallel program wasn't very large.

Scalability

- The word “scalable” has a wide variety of informal uses.
- A program is **scalable** if, by increasing the power of the system it’s run on (e.g., increasing the number of cores), we can obtain speedups over the program when it’s run on a less powerful system (e.g., a system with fewer cores).
- Suppose we run a parallel program with a fixed number of processes/threads and a fixed input size, and we obtain an efficiency E.
- Suppose we now increase the number of processes/threads that are used by the program. If we can find a corresponding rate of increase in the problem size so that the program always has efficiency E, then the program **is scalable**.
 - If we increase the number of processes/threads and keep the efficiency fixed without increasing problem size, the program is **strongly scalable**.
 - If we can keep the efficiency fixed by increasing the problem size at the same rate as we increase the number of processes/threads, then the program is **weakly scalable**.

Performance Tradeoff

- We have seen there are many factors that make it difficult to **write high performance parallel programs**
 - But it gets worse (more complex). . .
 - **Overheads, Amdahl's Law, idle time**, etc. reduce performance.
 - However, trying **to reduce the effect of one factor can increase the effect of another factor**.

Performance Tradeoff

- Three main trade-offs
 - Communication vs. computation
 - Memory vs. parallelism
 - Overhead vs. parallelism

Communication vs. Computation

- **Communication can add large overhead.** E.g., NUMA
- **Overlap communication and computation**
 - Identify communication that is independent of the computation.
 - Send communication message, and then do computation while waiting for message reply. Hide communication latency.
- **Can be reduced by extra computation**
 - Redundant/extr computation
 - Rather than requesting some data from memory, processor calculates data values by itself (if this is possible).
 - Objective: do extra computation, if it costs less than communication

Memory vs. Parallelism

- We saw that if we use **more memory, we can improve parallelism**
 - **Privatization**
 - Make **variables** to remove false dependencies
 - **Padding**
 - **Spread variables out in memory** (by allocating extra memory) to remove false sharing
 - Sometimes the wasted space can be used for other data, but makes code more complex.
- Improving memory efficiency can reduce parallelism.

Overhead vs. Parallelism

- Want to **minimize overhead**
- Increasing parallelism increases overhead
 - E.g., **create more threads to use more procs.**
 - But creating threads adds overhead
 - **More threads increase contention**
 - If the problem size is fixed then:
 - **Increase threads:** less work per thread.
 - **Overhead larger part** of each thread's time.
 - **Less chance** to hide communication cost.

Summary

- **Parallel software**
 - SPMD
 - Coordinating the processes/threads
 - Shared-memory
 - Distributed-memory
 - Input and output
- **Measuring performance**
 - Speedup = T_s / T_p
 - Efficiency = speedup / P
 - Amdahl's law
 - Scalability
 - Taking timings of MIMD programs

Summary

- **Sources of performance loss**
 - **Overheads**
 - Communication
 - Synchronization
 - Computation
 - Memory
 - **Non-parallelizable code**
 - **Contention**
 - **Idle time**
- **Performance trade-offs**
 - Communication vs. computation
 - Memory vs. parallelism
 - Overhead vs. parallelism